

Final Project, Digital Geometry Processing

Shayan Hoshyari
Student #: 81382153

December 2016

Introduction

In this project an adaptive surface remesher has been developed based on the paper [1]. An algorithm is proposed to increase the quality of surface manifold meshes using four fundamental operations, i.e., edge split, edge collapse, edge flip, and vertex relocation.

In the first section, strategies are introduced to perform each of the fundamental operations. Subsequently, the primitive operations are glued together in the form of a global algorithm. Finally, some results are presented.

Primitive Strategies

The concepts of edge collapse, edge flip, edge split, and vertex relocation are rather straightforward in two dimensions. However, generalizing them to manifold surfaces needs extra steps to be taken. In this section we will introduce these core operations. Namely controlling mesh fidelity, projection of points on the initial mesh, geodesic polar mapping, Delaunay edge flips, area based vertex relocation, and Laplacian smoothing.

Controlling Mesh Fidelity

In order to prevent deviating from the initial surface we require the normals of each triangle to be close to the normals at its vertices. Furthermore, the normals at the vertices should be close to one another as well. Mathematically, this can be expressed as:

$$\begin{aligned}\min(N_i, N_j) &> \cos(\theta_1) \quad i, j \in T \\ \min(N_i, N_T) &> \cos(\theta_2) \quad i \in T\end{aligned}$$

Where θ_1 and θ_2 are input parameters and are set to 20 degrees in this work. N_i , and N_j are the normals at any vertex of the triangle T , and N_T is the normal of the triangle itself. If any of our operations defy these criteria, we simply do not perform them.

Projection on the Initial Surface

All our operations that include vertex insertion or relocation express the location of a point via the pair (T, b) . Where T is a triangle on the current mesh to whom this new point belongs, and b is the barycentric coordinate of the point on T . One strategy is to choose the location of the point according to the current mesh, i.e., just interpolate a point on T using b . A more advanced strategy uses overlapping parameterized patches to project the point onto the original surface. Figure 1 shows how the former can deform the surface while the latter preserves the fidelity of the surface. Both strategies are implemented in this work.

The overlapping parameterized patches work in the following way. First three triangles on the original mesh that contain vertices of T are identified. Lets call them T_o . Then using the BFS search we create a patch containing all three members of T_o . The patch should be isomorphic to a disk so we prevent adding triangles that can create holes. Such a triangle is shown in Figure 2. Then the ears of the patch are trimmed and the patch is mapped to a unit disk using the CGAL[2] library. Then we interpolate the point on the mapped surface, and find the triangle to whom the interpolated point belongs, along with its barycentric coordinates. Lets call them (T_t, b_t) . Finally, we use T_t and b_t on the initial mesh to find the desired location of the point. Figure 3 shows several steps of the process.

Geodesic Polar Mapping

Delaunay edge flips, Laplacian smoothing, and area based vertex relocation will all be defined in two dimensions. In order to perform them on surface meshes we first map the neighborhood of a vertex or an edge to two dimension, and then perform the operation. In the case of vertex relocation we use the barycentric coordinate of the mapped triangle to map the point back to three dimensions.

Figure 4 shows the schematic of such a mapping. In the case of an edge one of the triangles is just rotated around the common edge until both are co-planar. In the case of a vertex we preserve the length of edges connecting the vertex to the neighboring ones but scale the angles by a constant factor so that they sum to 2π .

Delaunay Edge Flips

In order to ensure that all edges are Delaunay in the geodesic polar mapping sense, we use algorithm 1 taken from [4].

Algorithm 1 Delaunay Flips

- 1: Push all non-locally interior edges of T on stack and mark them.
 - 2: **while** stack is non-empty **do**
 - 3: $pq = \text{pop}()$
 - 4: unmark pq
 - 5: **if** pq is non-locally Delaunay **then**
 - 6: Replace by the edge connecting the respective third vertices of the two incident triangles
 - 7: Push other four edges of the two triangles into the stack if unmarked
 - 8: **end if**
 - 9: **end while**
-




Figure 1: Comparison of using the current mesh vs. the original mesh for vertex insertion and relocation. The left figure shows the more advanced overlapping parameterized patches technique which is clearly closer to the original surface than the right column which uses the current mesh for vertex relocation. The original surface is shown for comparison.




Figure 2: The starred triangle is an example that should not be added to the patch. The dashed triangles are already members of the patch. We identify the starred triangle by the fact that it only has one neighbor which is inside the patch and a vertex which belongs to triangles that are members of the patch.




Figure 3: Overlapping parameterized patches: The left figure shows three triangles T_o on the initial mesh that contain an imaginary triangle T_c on the current mesh. Using the BFS search a patch is created containing all these three triangles. In the second figure the ears of the patch are trimmed. The third figure shows an isolated view of the patch. In the fourth figure the patch is mapped to a unit disk.




Figure 4: Schematic of geodesic polar mapping of the neighborhood of an edge or a vertex. The picture is taken from reference [3].

Area Based Vertex Relocation

The area based vertex relocation ensures that the vertices have a uniform distribution. As noted by algorithm 2 this operation has to be followed by Delaunay edge flips. The idea is pretty simple, for each interior vertex in the mesh we displace it so that the area of its adjacent triangles are equalized. It reduces to solving a linear least squares problem in the following form:

$$\sum (A_i(x, y) - \mu_i \sum A_i)^2 = 0$$

Where A_i s are the area of the adjacent triangles, and (x, y) is the location of the vertex. We solve this problem using QR decomposition from the “Armadillo” linear algebra library. The parameters μ_i are chosen to be 1 in this work, but can generally be used to control the grading of the mesh.

Laplacian Smoothing

When a mesh is almost good, this operation can make it “very good”. For each vertex in the mesh we displace it to the average location of its neighbors in the geodesic polar mapped space.

Gluing It All Together

Having all the small pieces working, we will combine them to create our remesher. First we read a target number of vertices from the command line switches. If the number of target vertices is more than the current vertices we start splitting the biggest edge of bad triangles. In each round of insertion we sort the edges according to the quality of adjacent triangles, and start splitting the ones with the worst neighboring triangles. We do not split an edge when any of its adjacent triangles have been split in the current round, neither when it is not the biggest edge of any of the adjacent triangles.

If the number of target vertices is less than the current number of vertices, we do several rounds of edge collapse. Again we sort edges according to the quality of neighboring triangles. This time we collapse edges that have higher quality neighbors first. The edge being collapsed must not be too big compared to the other edges of the adjacent triangles, neither should any of those triangles have been collapsed in the same round.

After each round of insertion or edge collapse we do three rounds of area based vertex relocation followed by Delaunay edge flips.

Our insertion and collapse algorithm is rather naive, especially in the vertex reduction scheme where we do not do any insertion at all. A better approach would be to have all triangles in a priority queue, and smartly do a combination of collapse and insertion to the bad triangles.

After the desired number vertices have been reached we have the option to split all edges that are facing obtuse angles. This step should only be done once, as recursive performance of this action can create really tiny triangles that can cause numerical problems.

The next step is to perform ten rounds of area based remeshing. In each round all vertices will be relocated using the area based method three times followed by Delaunay edge flips. The quality of the mesh is further enhanced by performing Laplacian smoothing. Algorithm 2 concisely reviews all the involved steps.

Algorithm 2 Gluing the primitive operations together

- 1: **while** Target # of vertices is reached **do**
 - 2: Sort the edges according to ascending/descending adjacent triangle quality.
 - 3: Split/collapse the edges in the mentioned order.
 - 4: Do not collapse or split edges that share an adjacent triangle.
 - 5: Perform 3 rounds of area based vertex relocation.
 - 6: Perform Delaunay edge flips.
 - 7: **end while**
 - 8: Optionally split all edges facing obtuse angles (not recursively).
 - 9: Do the following 10 times: 3 rounds of area based vertex relocation followed by Delaunay edge flips.
 - 10: Perform 10 rounds of Laplacian smoothing.
-

Results

The performance of the code is evaluated by remeshing five different geometries. In the first four cases we use insertion to increase the number of vertices. The selected geometries are a sphere, the cow, Olliver’s Hand, and Nicolo’s face, taken from the aim@shape repository. We insert points via two different criteria, i.e., splitting the biggest edge of bad triangles, and removing obtuse angles. In the fifth case we remesh the horse model using edge collapse. To show the superiority of the overlapping parameterized patches technique, we also show the results of the current mesh interpolation for the horse case.

Figure 5 shows the initial and remeshed models for insertion cases. Both the insertion for removing obtuse angles, and splitting biggest edge of bad triangles work pretty well. The flat shaded remeshed surfaces bear the same pattern as the original ones, which shows how the method preserves surface fidelity, specifically in the hand and face geometries. There are some areas like the nose of the face model that have high initial fidelity error. These areas cannot be modified properly and are still poor even after remeshing. A workaround is to smooth out these areas with adaptive subdivision or PN triangles before starting the remeshing process.

Figure 6 shows the horse model which is remeshed via edge collapse. Although our vertex reduction scheme is rather naive and only collapses edges without insertion, the increase in mesh quality is still noticeable. We also see how the overlapping parameterized patches preserve the mesh fidelity, while using the current mesh to find vertex locations causes the mesh to shrink. The overlapping patches have only increased the run time by less than 50%, which shows they are as efficient as they are effective.

Table 1: Remesher performance. Run time and the percentage of triangles with different minimum angles are shown. The horse (a) and (b) cases use overlapping patches while horse (c) and (d) use interpolation on the current mesh.

Name	Vertex #	Run Time	Percentage of triangles with certain minimum angles					
			0 – 10	10 – 20	20 – 30	30 – 40	40 – 50	50 – 60
Sphere 0	482	—	0.00	13.33	20.00	26.67	40.00	0.00
Sphere a	930	0.77 sec	0.00	0.00	0.00	0.00	16.54	83.56
Sphere b	10000	6.76 sec	0.00	13.33	20.00	0.03	21.95	78.03
Cow 0	2904	—	3.76	15.23	29.91	30.96	16.25	3.89
Cow a	6000	4.28 sec	0.47	3.91	9.29	15.54	31.95	38.85
Cow b	20000	14.50 sec	0.18	1.49	5.15	14.87	31.03	47.28
Hand 0	2531	—	0.68	8.94	23.40	31.36	25.66	9.90
Hand a	4320	2.63 sec	0.17	0.74	2.11	5.86	31.35	59.88
Hand b	20000	15.20 sec	0.01	0.17	1.21	5.28	27.04	66.29
Face 0	2554	—	0.76	5.56	18.74	32.11	30.23	12.60
Face a	4129	2.53 sec	0.26	1.61	4.06	8.27	30.69	55.11
Face b	20000	14.70 sec	0.03	0.46	2.26	8.82	27.43	60.99
Horse 0	19851	—	0.02	0.44	6.37	29.47	43.00	20.71
Horse a	10000	19.63 sec	0.03	0.38	1.40	4.17	28.23	65.79
Horse b	5000	16.02 sec	0.09	1.18	4.64	10.59	31.73	51.77
Horse c	10000	14.60 sec	0.04	0.28	0.77	1.68	25.50	71.75
Horse d	50000	10.11 sec	0.14	0.92	2.31	5.00	28.68	62.95

Table 1 shows a summary of the remeshing performance for each case. The run time of the code is reasonable even up to 20K vertices, while the percentage of triangles with good minimum angles increases with inserting further vertices. The vertex reduction algorithm does not perform as good as the previous cases, as it only does edge collapse naively without any insertion.

Conclusion

In this project an adaptive remesher was developed based on the strategies introduced in the paper [1]. Primitive operations such as edge split, edge collapse, edge flip, and vertex relocation were performed to increase the quality of the manifold surface. As most of the primitive operations are defined in two-dimensions we used the geodesic polar mapping to map the adjacency of each vertex or edge into two-dimensions. To ensure the fidelity of the final surface, two fidelity error criteria were introduced based on the vertex and triangle normals. Any operation that defies them should not be performed. Furthermore, the parameterization of overlapping patches was used in order to project the relocated or inserted vertices to the original surface. Finally, in order to evaluate the performance of the code several models were remeshed and the results were presented.

There are several ways to improve the performance of the code. Firstly, the initial areas that have high fidelity errors can be smoothed out before remeshing the surface. Similarly, curved surfaces such as PN triangles can be used for presenting the initial surface in order to have smoother remeshed results. Another potential improvement would be to do edge split and collapse in a smarter way, such as keeping a priority queue and a criterion to decide whether to collapse or to split an edge of the worst triangle. Furthermore, feature edges can be used where the normals are duplicated on each feature vertex. This allows remeshing of faceted surfaces, such as a cylinder or a cube. Curvature sensitive triangle size grading can also improve the quality of the final meshes considerably.




Figure 5: The models remeshed via insertion. Left column is the original model. In the middle column insertion is only performed to remove the obtuse angles. In the last column target number of vertices is set to 20000.

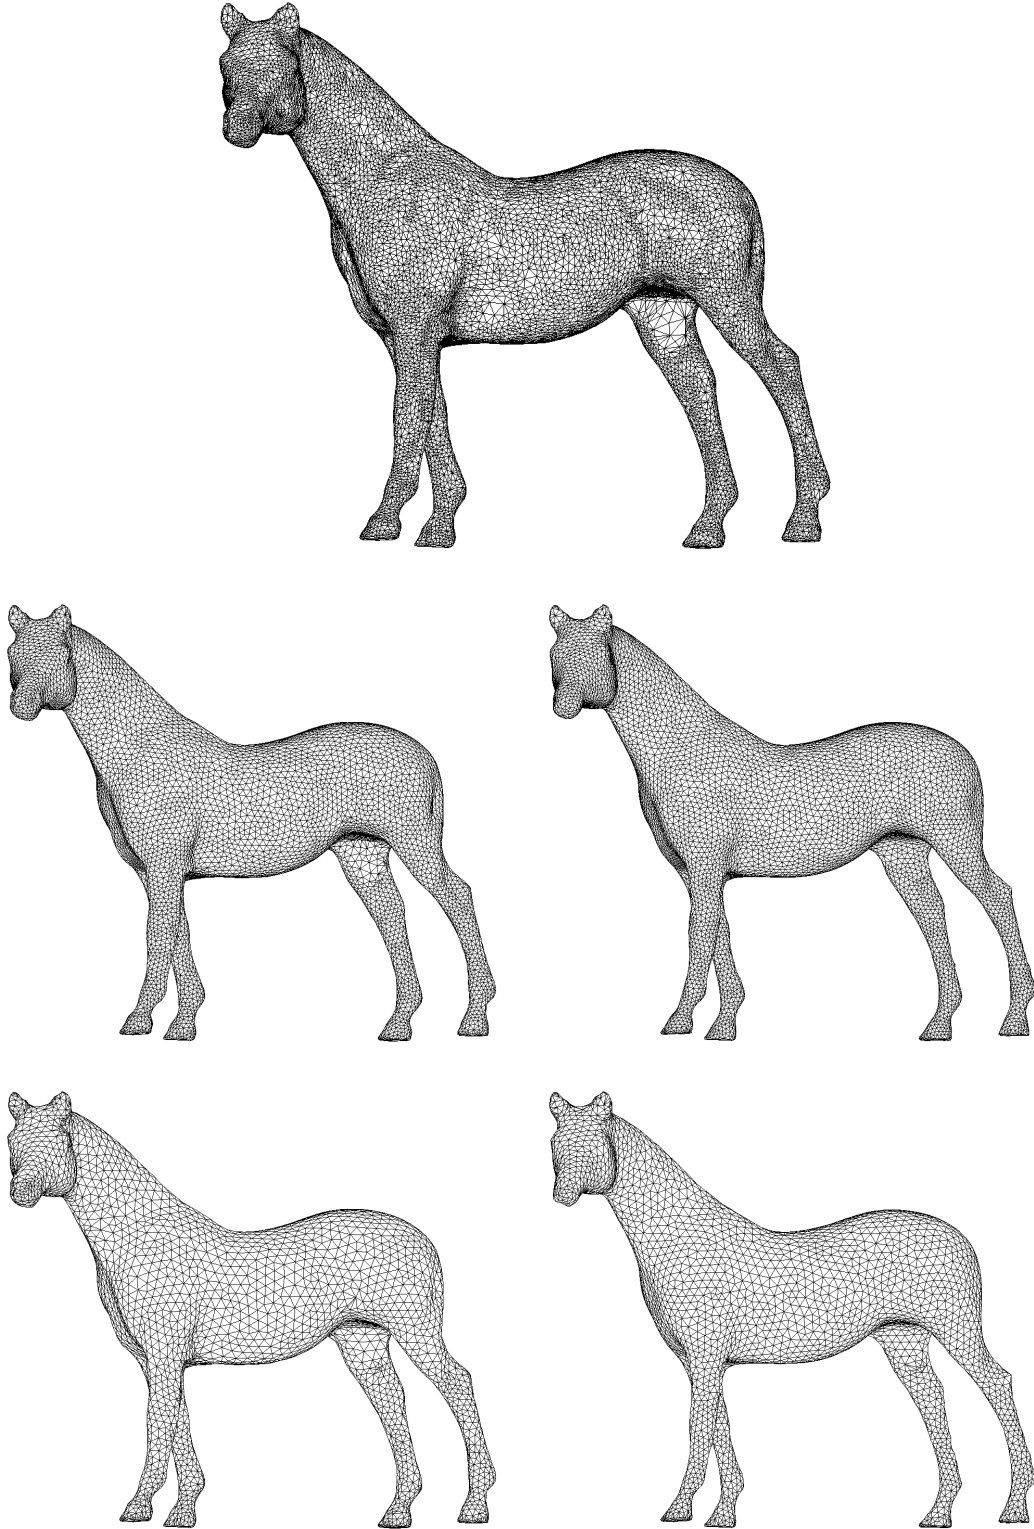


Figure 6: The horse model remeshed via edge collapse. The top figure shows the initial mesh. In the left column the new points are inserted on the initial surface via the overlapping parameterized patches while in the right column the new points are inserted on the current mesh. It is clear that the meshes on the right column have all shrunk, especially in the mouth and leg areas.

References

- [1] V. Surazhsky and C. Gotsman, “Explicit surface remeshing,” in *Proceedings of the 2003 Eurographics/ACM SIGGRAPH symposium on Geometry processing*, pp. 20–30, Eurographics Association, 2003.
- [2] L. Saboret, P. Alliez, and B. Levy, “Planar parameterization of triangulated surface meshes.” http://doc.cgal.org/latest/Surface_mesh_parameterization/index.html. Accessed: 2016-12-16.
- [3] E. L. Melv  r and M. Reimers, “Geodesic polar coordinates on polygonal meshes,” in *Computer Graphics Forum*, vol. 31, pp. 2423–2435, Wiley Online Library, 2012.
- [4] A. Ungor, “Edge flip algorithm for delaunay triangulation.” <https://www.cise.ufl.edu/~ungor/delaunay/delaunay/node5.html>. Accessed: 2016-12-16.